

Priya Nair

AI Engineer

priya.nair@email.com | (555) 908-3345 | Seattle, WA | linkedin.com/in/priyanair-ai | github.com/priyanair

SUMMARY

AI Engineer with 6 years of experience designing and deploying machine learning and large language model systems in production. Specializes in retrieval-augmented generation, model fine-tuning, and building scalable ML pipelines. Skilled at bridging research and engineering to ship reliable, cost-efficient AI products.

CORE SKILLS

Languages	Python, SQL, Bash
ML/DL Frameworks	PyTorch, TensorFlow, Hugging Face Transformers, LangChain, LlamaIndex
LLM Systems	RAG pipelines, prompt engineering, fine-tuning (LoRA/QLoRA), vector search
MLOps	MLflow, Kubeflow, Docker, Kubernetes, Airflow, CI/CD for ML
Cloud	AWS (SageMaker, Bedrock), GCP (Vertex AI), Azure ML
Data	PostgreSQL, Pinecone, Weaviate, Spark, Pandas

PROFESSIONAL EXPERIENCE

AI Engineer — Cognitive Labs

Feb 2022 – Present | Seattle, WA

- Built and shipped a retrieval-augmented generation (RAG) system powering an enterprise support assistant, reducing average ticket resolution time by 38%.
- Fine-tuned open-source LLMs (7B–13B parameters) using LoRA on domain-specific data, improving task accuracy by 24% over base models.
- Designed a vector search pipeline using Pinecone and custom embedding models, achieving sub-200ms retrieval latency at scale.
- Deployed model-serving infrastructure on AWS SageMaker with autoscaling, cutting inference costs by 30% versus prior setup.
- Partnered with product teams to define evaluation metrics and guardrails, reducing hallucination-related incidents by 45%.

Machine Learning Engineer — DataSphere Analytics

Aug 2019 – Jan 2022 | Portland, OR

- Developed and productionized classification and forecasting models in PyTorch for demand-planning clients, improving forecast accuracy by 18%.
- Built automated training and evaluation pipelines using Airflow and MLflow, reducing model iteration time from days to hours.
- Implemented feature stores and data validation checks that cut data-quality-related model failures by 50%.
- Presented model performance and business impact analyses to non-technical stakeholders on a quarterly basis.

Data Scientist — Insight Metrics

Jul 2017 – Jul 2019 | Portland, OR

- Built predictive models for customer churn using scikit-learn, contributing to a 12% reduction in churn rate.
- Conducted A/B testing and statistical analysis to guide product and marketing decisions.
- Automated reporting workflows in Python, saving the analytics team roughly 10 hours per week.

EDUCATION

M.S. in Computer Science, Machine Learning Track — University of Washington

Graduated 2017

B.S. in Computer Engineering — Oregon State University

Graduated 2015

CERTIFICATIONS

- AWS Certified Machine Learning – Specialty
- DeepLearning.AI Generative AI with LLMs Specialization